
Dawn Documentation

Release 0.0.1

Fabian Thuring, Carlos Osuna and Tobias Wicky

ETH Zurich and MeteoSwiss

October 04, 2017

THE BASICS

1	About this project	1
2	First steps	3
3	Stencil Intermediate Representation	7
4	Dawn toolchain	9
5	ISIR	11
6	GridTools	13
7	Dawn Coding standard	15
8	CMake Modules	17

ABOUT THIS PROJECT

Dawn is a compiler toolchain for developing high-level DSLs for geophysical fluid dynamics models. Development productivity of large scientific codes, like geophysical fluid dynamics (GFD) models, decreased drastically in recent times due to the fact these community models often have to run efficiently on multiple computing architectures that impose the use of different programming models. Embedded domain specific languages (EDSLs) in C++ have been used successfully in the past to enable a clear separation of concerns between the domain algorithms and the implementation strategy, thus allowing a single source code which is performance portable across heterogeneous architectures. Yet, using EDSLs in a correct and efficient manner often requires expert knowledge in high-performance computing. In addition, current DSLs are usually designed and developed for a specific scientific model with little to no reusability among DSLs.

We introduce a new compiler framework, Dawn, that decouples optimization and code generation from high level DSLs. By exposing a Stencil Intermediate Representation (SIR), we allow users to share the toolchain to optimize and generate code among several DSLs. This allows the design of expressive, concise DSLs that can focus on applicability and don't need to bother with the aspect of high-performance computing. Using Dawn as a common compiler infrastructure can drastically reduce development and maintenance effort, while increasing the performance of the generated code, for new and existing DSLs in the GFD model domain.

Core Features

- Dawn allows the user to generate fast performing code for several back-ends from a relatively simple Stencil Intermediate Representation (SIR).
- Dawn exposes several APIs in different languages (C++, Java, Python) to parse and process the SIR.
- Dawn is able to generate code to be run on Distributed Memory Machines based on MPI, Machines with access to GPUs based on CUDA (via the gridtools EDSL) as well as naive C++ code with close to no parallelism for debugging.
- Dawn offers a wide range of optimization and static analysis passes to guarantee correctness as well as performance of the generated parallel program.

FIRST STEPS

This section demonstrates the basic features of Dawn. Before getting started, make sure that development environment is set up to compile the included set of test cases.

Compiling the library

Linux

On Linux, you'll need to install a C++11 toolchain (e.g **gcc** or **clang**) as well as **CMake** (≥ 3.3). For example, on Ubuntu the following will do

```
sudo apt-get install g++ cmake
```

After installing the prerequisites, run

```
mkdir build
cd build
cmake ..
make install -j 4
```

The last line will both install and compile the library locally in **<dawn-dir>/install/**.

Mac OS

On Mac OS, you'll need to install Xcode as well as **CMake** (≥ 3.3). First, make sure you have the Xcode Command Line Tools installed

```
xcode-select --install
```

If you are using **Homebrew**, use the following to install **CMake**:

```
brew update
brew install cmake
```

After installing the prerequisites, run

```
mkdir build
cd build
cmake ..
make install -j 4
```

The last line will both install and compile the library locally in **<dawn-dir>/install/**.

Linking the library

We suggest using [CMake](#) for a smooth integration with your project. The installation of Dawn provides a find configuration script (which can be used with `find_package`) that should be installed into your system.

In your projects `CMakeLists.txt`, simply add

```
find_package(Dawn)
```

to import all the necessary information.

Note: The script is located in `<dawn-install-dir>/cmake`. If CMake has trouble finding the script you may pass this location directly to CMake by adding `-DDawn_DIR=<dawn-install-dir>/cmake` to the command line flags.

If Dawn is found, the following CMake variables will be defined

Variable	Explanation
DAWN_FOUND	True if headers and libraries of Dawn were found.
DAWN_ROOT	Installation prefix of Dawn.
DAWN_VERSION	Version of Dawn (format X.Y.Z).
DAWN_ASSERTS	True if Dawn was compiled with asserts.
DAWN_INCLUDE_DIRS	Dawn include directories.
DAWN_LIBRARY_DIRS	Link directories for Dawn libraries.
DAWN_LIBRARY	Library to link against (this is an alias of <code>DAWN_STATIC_LIBRARY</code>).
DAWN_STATIC_LIBRARY	Static library of Dawn.
DAWN_HAS_SHARED_LIBRARY	True if the shared library of Dawn is available.
DAWN_SHARED_LIBRARY	Shared library of Dawn.

Finally, just link the static library of Dawn to your own library or executable

```
target_link_libraries(${target} ... PUBLIC Dawn::DawnStatic)
```

or use `Dawn::DawnShared` instead of `Dawn::DawnStatic` if you want to linkt against the shared library.

Example

If we want to link our file (`foo.cpp`) against the static library of Dawn, we may use something like

```
# Find the Dawn library (abort if we cannot find it!)
find_package(Dawn REQUIRED)

# Expose the Dawn include directories
include_directories(SYSTEM ${DAWN_INCLUDE_DIRS})

# Compile foo and link against Dawn
add_executable(foo foo.cpp)
target_link_libraries(foo PUBLIC Dawn::DawnStatic)

# Dawn requires atleast C++11 (you can also set this globally!)
set_property(TARGET foo PROPERTY CXX_STANDARD 11)
```

Using the library

TODO:


```
#include <dawn/Dawn.h>

int main() {

    return 0;
}
```


STENCIL INTERMEDIATE REPRESENTATION

Developing a single DSL that is able to support all numerical methods and computational patterns present in GFD models present a serious challenge due to the wide domain it needs to cover. In this project, we accept the reality that multiple scientific communities desire to develop their own DSL language, tailored for their model and needs. The Stencil Intermediate Representation (SIR) allows to define multiple high level DSLs in a lightweight manner by reusing most of the complex toolchain i.e the Dawn library. In addition, a standardized SIR has another major advantage: It allows to easily interact with third-party developers and hardware manufacturers. For instance, the SIR of the [COSMO](#) atmospheric model, serialized to a mark-up Language like [JSON](#), can be distributed to hardware vendors, which in turn have their proprietary, in-house compilers based on the SIR and can return plain C/C++ or CUDA code. This frees third-party developers from compiling the models and hopefully improves collaboration.

In this section we define the specification of the SIR.

DAWN TOOLCHAIN

The toolchain of Dawn contains several modules and is organized in passes, allowing to customize the process by enabling, disabling or change the order of execution of the individual passes. The default order of execution, depicted in Figure 4.1, starts by translating the SIR into an *internal* SIR (*ISIR*) which enriches the SIR with various information, like the computation of the access patterns, required by the following passes. As the initial construction of the ISIR generates a parallel program that is most likely incorrect under the given parallel model, we invoke several passes (*Parallel Model Builder*) to build a valid parallel program. The *Safety Checker* passes will run various static analysis routines and the *Optimizer* will perform several optimizations such as reordering and fusing groups of statements, setting software managed caches or reducing shared-memory pressure. Finally, the code generation will translate the ISIR into the main DSL of GridTools or directly to C++.




Fig. 4.1: Default order of execution of the passes grouped in *Parallel Model Builder* (orange), *Safety Checker* (yellow) and *Optimizer* (green).

Parallel Model Builder

The parallel model builder passes transform the initial Stage and Multi-Stage arrangement into a valid parallel program under the *GridTools parallel model*.

Safety Checker

TODO

Optimizer

TODO

TODO

GRIDTOOLS

TODO

Parallel Model

TODO

DAWN CODING STANDARD

Introduction

This document attempts to describe a few coding standards that are being used in the Dawn source tree. Although no coding standards should be regarded as absolute requirements to be followed in all instances, you should try to follow the [LLVM](#) coding standard. We deviate from the LLVM standard in certain areas as listed in the following.

Supported C++11 Language and Library Features

While LLVM restricts the usage of C++11, we do not impose any constraints on the allowed features. While we generally allow the usage of C++ exceptions, you should avoid them at the interface boundary to make exposing the API more convenient - [RTTI](#) should be avoided.

Style Issues

Naming of Types, Functions, Variables, and Enumerators

- **Variable Names** should start with a lower-case letter (e.g `textFileReader`) as opposed to LLVM which starts with an uppercase-letter.

For types, functions and enumerators you should follow the LLVM style.

Source Code Formatting

Source Code Width

You may use up to 100 columns of text when writing your code (instead of 80 as dictated by LLVM). The rationale behind this is that we do not live in the 90s any more and people usually have wide screens.

Spaces Before Parentheses

You should never use spaces before parentheses. For example, this is good

```
if(x) ...
for(i = 0; i != 100; ++i) ...

somefunc(42);
assert(3 != 4 && "laws of math are failing me");

A = foo(42, 92) + bar(X);
```

and this is bad

```
if (x) ...
for (i = 0; i != 100; ++i) ...

somefunc (42);
assert (3 != 4 && "laws of math are failing me");

A = foo (42, 92) + bar (X);
```

Pointer and Reference Alignment

You should always align pointers and references on the left i.e directly following the type. For example, this is considered good

```
int* a = ...
int& b = ...

const char** ptr = ...
```

while the following is considered bad

```
int *a = ...
int &b = ...

const char **ptr = ...
```

Clang Format

To enforce most of these coding standards, CMake can be configured to run `clang-format` on each file with

```
make format
```

The clang-format file is located in the root directory at `<dawn-dir>/clangformat`.

CMAKE MODULES

This section describes the CMake functionality of Dawn and its subprojects.

Installation

Add the `<dawn>/cmake/modules` directory to the `CMAKE_MODULE_PATH` to use the functions, macros and modules:

```
list(APPEND CMAKE_MODULE_PATH "<dawn>/cmake/modules")
```

Note that all Dawn projects contain a `dawn_cmake_init` macro which tries to find the CMake modules of Dawn.

```
include(DawnCMakeInit)
dawn_cmake_init()
```

Functions & Macros

Each function and macro uses a `snake-case` identifier and is defined in a separate file using the corresponding `camel-case` filename. For example, to use the function `dawn_add_target_clean_all` include the file `DawnAddTargetCleanAll`.

```
include(DawnAddTargetCleanAll)
dawn_add_target_clean_all()
```

dawn_get_compiler_info

Get the identification string of the compiler.

```
dawn_get_compiler()
```

The function defines the following variable:

DAWN_COMPILER_STRING Variable which will contain the identification string of the compiler.

and conditionally one of the following:

DAWN_COMPILER_MSVC Set to 1 if the compiler in use is MSVC

DAWN_COMPILER_GNU Set to 1 if the compiler in use is GNU

DAWN_COMPILER_CLANG Set to 1 if the compiler in use is Clang

dawn_add_executable

Compile the given objects into a runnable executable (.exe).

```
dawn_add_executable(NAME SOURCES DEPENDS [OUTPUT_DIR])
```

NAME Name of the executable as well as the CMake target to build it.

SOURCES List of source files making up the executable.

DEPENDS List of external libraries and/or CMake targets to link against.

OUTPUT_DIR [optional] Directory to place the executable (e.g. \${CMAKE_BINARY_DIR}/bin).

dawn_add_library

Compile the given sources into an archive or object library. This will provide the CMake targets <NAME>Objects, <NAME>Static and <NAME>Shared respectively.

```
dawn_add_library(NAME SOURCES DEPENDS [OUTPUT_DIR])
```

NAME Name of the executable as well as the CMake target to build it.

SOURCES List of source files making up the executable.

DEPENDS List of external libraries and/or CMake targets to link against.

OBJECT [optional] Create a CMake object library.

ARCHIVE [optional] Create a static and, if BUILD_SHARED_LIBS is ON, shared archive.

dawn_add_target_clang_format

Provide a `format` target which runs `clang-format` recursively on all files in the provided directories.

```
dawn_add_target_clang_format(DIRECTORIES PATTERN)
```

DIRECTORIES Directories to recursively traverse to find all files with extensions matching `EXTENSION`.

EXTENSION Extension to match, separated by `;`. For example: `".h; .cpp"`.

dawn_add_target_clean_all

Provide a `clean-all` target which clears the CMake cache and all related CMake files and directories. This effectively removes the following files/directories:

- \${CMAKE_BINARY_DIR}/CMakeCache.txt
- \${CMAKE_BINARY_DIR}/CTestTestfile.cmake
- \${CMAKE_BINARY_DIR}/cmake_install.cmake
- \${CMAKE_BINARY_DIR}/CMakeFiles

```
dawn_add_target_clean_all([ARGN...])
```

ARGN Additional files or directories to remove.

dawn_add_unittest

Compile the given objects into a runnable unittest executable (.exe) and register it within CTest. Note that to enable CTest you need to call the builtin command `enable_testing()` in the source root.

```
dawn_add_unittest(NAME SOURCES DEPENDS [OUTPUT_DIR GTEST_ARGS])
```

NAME Name of the unittest executable as well as the CMake target to build it.

SOURCES List of source files making up the executable.

DEPENDS List of external libraries and/or CMake targets to link against.

OUTPUT_DIR [optional] Directory to place the executable (e.g. `${CMAKE_BINARY_DIR}/bin`).

GTEST_ARGS [optional] Arguments passed to the created GTest executable (e.g. `--gtest_color=yes`)

dawn_check_and_set_cxx_flag

Test if the C++ compiler flag is supported and if so, add it to the `CMAKE_CXX_FLAGS`.

```
dawn_check_and_set_cxx_flag(FLAG NAME)
```

FLAG Compiler flag to check (e.g. `-O3`).

NAME Name of the check (e.g. `HAVE_GCC_O3`).

dawn_check_cxx_flag

Test if the C++ compiler flag is supported.

```
dawn_check_cxx_flag(FLAG NAME)
```

FLAG Compiler flag to check (e.g. `-O3`).

NAME Name of the check (e.g. `HAVE_GCC_O3`).

dawn_check_in_source_build

Abort with a meaningful error message if `CMAKE_BINARY_DIR` matches `CMAKE_SOURCE_DIR`.

```
dawn_check_in_source_build()
```

dawn_combine_libraries

Combine multiple object libraries to a single static and, if `BUILD_SHARED_LIBS` is ON, shared library. The CMake target of the library is `<NAME>Static` and `<NAME>Shared` respectively. This will also provide an install target for the libraries as well as an export via `<NAME>Targets`.

```
dawn_combine_libraries(NAME OBJECTS DEPENDS)
```

NAME Name of the library.

OBJECTS Object libraries to combine (see *dawn_add_library*).

INSTALL_DESTINATION Destination (relative to `CMAKE_INSTALL_PREFIX`) to install the libraries.

DEPENDS [optional] List of external libraries and/or CMake targets treated as dependencies of the library.

dawn_configure_file

Configure a file and store the output in the same directory hierarchy as the input file while substituting `CMAKE_SOURCE_DIR` for `CMAKE_BINARY_DIR`. If the input file has a `.cmake` extension, it will be removed in the output. This will also handle the installation of the file correctly by replacing `src/` with `include/`.

```
dawn_configure_file(FILE)
```

FILE File to configure.

Example

```
dawn_configure_file(${CMAKE_SOURCE_DIR}/src/dawn/Config.h.cmake)
```

This stores the configured file in `${CMAKE_BINARY_DIR}/src/dawn/Config.h` and, in addition, it will install it in `${CMAKE_INSTALL_PREFIX}/include/dawn/Config.h` (hence it will strip "src/").

dawn_create_package_string

Create a package string including the location of the library/include directory and version string. Note that the package has to be exported via *dawn_export_package*.

```
dawn_create_package_string(NAME PACKAGE_STRING [NOT_FOUND_STRING])
```

NAME Name of the package.

PACKAGE_STRING Returned package string.

NOT_FOUND_STRING String to print if the package was not found, defaults to "Not FOUND" (optional).

Example

```
dawn_export_package(Foo FOUND TRUE LIBRARIES "/usr/lib/libfoo.so" VERSION "1.2.3")
dawn_create_package_string(Foo FooPackageStr)
message(${FooPackageStr})
```

dawn_export_package

Export a package by defining variable for its libraries, include directories, definition and version.

```
dawn_export_package(NAME FOUND [LIBRARIES] [INCLUDE_DIRS]
                   [DEFINITIONS] [VERSION] [EXECUTABLE])
```

NAME Name of the package.

FOUND True if the package was found.

LIBRARIES List of libraries to link against (optional).

INCLUDE_DIRS List of include directories required by this package (optional).

DEFINITIONS List of definitions required by the package (optional).

EXECUTABLE Executable of the package (optional).

VERSION Version of the package (optional).

The following variables are defined:

DAWN_<NAME>_FOUND True if package was found.

DAWN_<NAME>_LIBRARIES Libraries of the package to link against.

DAWN_<NAME>_INCLUDE_DIRS Include directories required by this package.

DAWN_<NAME>_DEFINITIONS Definitions required by the package.

DAWN_<NAME>_EXECUTABLE Executable of the package.

DAWN_<NAME>_VERSION Version string of the package.

To create a formatted string of the exported packages *dawn_create_package_string*.

dawn_find_python_module

Try to find a specific python module.

```
dawn_find_python_module(MODULE [REQUIRED])
```

MODULE Python module to find.

REQUIRED If set, abort with an error if MODULE was not found.

The module makes use of the PYTHON_EXECUTABLE which needs to be set before calling this function (see [FindPythonInterp](#)).

Variables defined:

PYTHON_<module>_FOUND Python module is available where <module> is the python module to search for in all uppercase.

PYTHON_<module>_PATH Path to the module where <module> is the python module to search for in all uppercase.

dawn_get_architecture_info

Get the identification of the architecture.

```
dawn_get_architecture_info()
```

The function defines the following variable:

DAWN_ARCHITECTURE_STRING String of the architecture.

dawn_get_compiler_info

Get the identification string of the compiler.

```
dawn_get_compiler()
```

The function defines the following variable:

DAWN_COMPILER_STRING String of the currently used compiler.

and conditionally one of the following:

DAWN_COMPILER_MSVC Set to 1 if the compiler in use is MSVC

DAWN_COMPILER_GNU Set to 1 if the compiler in use is GNU

DAWN_COMPILER_CLANG Set to 1 if the compiler in use is Clang

dawn_get_git_head_revision

These functions force a re-configure on each git commit so that you can trust the values of the variables in your build system. It returns the refspec and sha1 hash of the current head revision.

```
dawn_get_git_head_revision(HEAD_REF HEAD_HASH)
```

HEAD_REF The Refspec.

HEAD_HASH sha1 hash of the current head revision.

dawn_get_platform_info

Get the identification string of the platform.

```
dawn_get_platform_info()
```

The functions defines the following variable:

DAWN_PLATFORM_STRING String of the platform.

and conditionally the following:

DAWN_ON_WIN32 Set to 1 if the platform is Win32-ish

DAWN_ON_UNIX Set to 1 if the platform is Unix-ish

DAWN_ON_APPLE Set to 1 if the platform is Darwin

DAWN_ON_LINUX Set to 1 if the platform is Linux

dawn_get_script_dir

Get the directory of the scripts located <dawn-root>/cmake/scripts.

```
dawn_get_script_dir (SCRIPT_DIR_VAR)
```

SCRIPT_DIR_VAR Variable which will contain the script directory on output.

Note: This function is for internal use only.

dawn_include_guard

Prevent frequently-included CMake files from being re-parsed multiple times.

```
dawn_include_guard()
```

dawn_make_string_pair

Create a formatted string of the given string pair.

```
dawn_make_string_pair(FIRST SECOND NUM_CHARS OUT_STRING)
```

FIRST First string.

SECOND Second string.

NUM_CHARS Characters before **SECOND** is printed.

OUT_STRING Output string.

This is equivalent to the python snippet:

```
OUT_STRING = "%s%s%s" % (FIRST, " " * (NUM_CHARS - len(FIRST)), SECOND)
```

Example

```
dawn_make_string_pair(Foo Bar1 10 out_var1)
dawn_make_string_pair(LongerFoo Bar2 10 out_var2)
message(${out_var1}) # "Foo      Bar1"
message(${out_var2}) # "LongerFoo Bar2"
```

dawn_report_result

Report a list of strings, encompassed by a box.

```
dawn_report_result(HEADER [ARGN...])
```

HEADER Header to print

ARGN List of strings to report

dawn_set_cxx_standard

Set the minimum standard of C++.

```
dawn_set_cxx_standard(MIN_CXX_STANDARD)
```

MIN_CXX_STANDARD Minimum C++ standard which needs to be supported, one of [c++11, c++14, c++1z, c++17, c++2a].

Modules

Load settings for an external project via `find_package`.

FindClang

```
find_package(Clang [REQUIRED] [QUIET])
```

Find the `Clang` libraries and includes. Uses the same include and library paths detected by *FindLLVM*.

Result Variables

Defines the following variables:

CLANG_FOUND True if Clang was found.

CLANG_INCLUDE_DIRS Where to find Clang includes.

CLANG_LIBS List of all Clang libraries.

CLANG_LIBCLANG_LIB C Interface to Clang.

CLANG_LIBRARY_DIRS Where to find Clang libraries.

CLANG_RESSOURCE_INCLUDE_PATH Path to the internal Clang headers needed by Frontend tools.

CLANG_<COMPONENT>_LIB Path to the <COMPONENT> library.

Hints

Set `LLVM_ROOT` to a directory that contains a LLVM/Clang installation.

FindLLVM

```
find_package(LLVM [REQUIRED] [QUIET])
```

Find the LLVM libraries and includes.

Result Variables

Defines the following variables:

LLVM_FOUND True if llvm found.

LLVM_INCLUDE_DIRS Where to find llvm include files

LLVM_LIBRARY_DIRS Where to find llvm libs

LLVM_CPPFLAGS LLVM preprocessor flags

LLVM_LDFLAGS LLVM linker flags.

LLVM_CXXFLAGS LLVM c++ compiler flags

LLVM_INSTALL_PREFIX LLVM installation prefix

LLVM_VERSION Version string (`llvm-config --version`)

Hints

Set `LLVM_ROOT` to a directory that contains a LLVM installation.

FindSphinx

```
find_package(Sphinx [REQUIRED] [QUIET])
```

This module locates `Sphinx`.

Result Variables

Defines the following variables:

SPHINX_FOUND System has sphinx.

SPHINX_EXECUTABLE Path to sphinx-build executable.

Hints

You can directly set `SPHINX_EXECUTABLE` if the module has trouble finding sphinx.

Findbash

```
find_package(bash [REQUIRED] [QUIET])
```

This module locates `bash`. `BASH_FOUND` will report if bash has been found.

Result Variables

Defines the following variables:

BASH_FOUND System has bash.

BASH_EXECUTABLE Path to the bash executable.

Hints

You can directly set `BASH_EXECUTABLE` if the module has trouble finding bash.

Findccache

```
find_package(ccache [REQUIRED] [QUIET])
```

This module locates `ccache`. `CCACHE_FOUND` will report if ccache has been found.

Result Variables

Defines the following variables:

CCACHE_FOUND System has ccache.

CCACHE_EXECUTABLE Path to the ccache executable.

Hints

You can directly set `CCACHE_EXECUTABLE` if the module has trouble finding ccache.

Findclang-format

```
find_package(clang-format [REQUIRED] [QUIET])
```

This module locates `clang-format`. `CLANG-FORMAT_FOUND` will report if clang-format has been found.

Result Variables

Defines the following variables:

CLANG-FORMAT_FOUND System has clang-format.

CLANG-FORMAT_EXECUTABLE Path to the clang-format executable.

Hints

You can directly set `CLANG-FORMAT_EXECUTABLE` if the module has trouble finding clang-format.